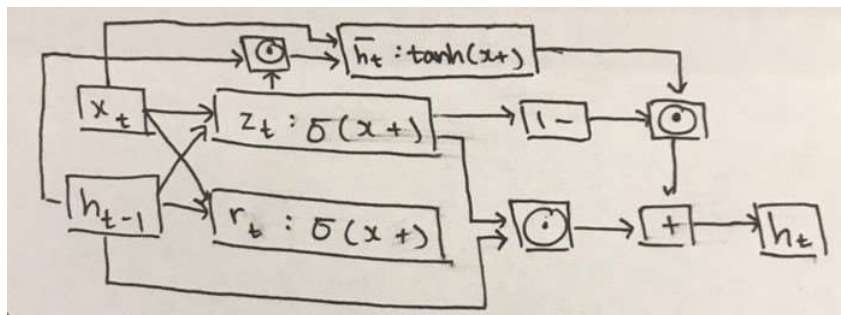


1.1:


$$\begin{aligned} W^z &: d_h \times d_x \\ W^r &: d_h \times d_x \\ W &: d_h \times d_x \\ U^z &: d_h \times d_h \\ U^r &: d_h \times d_h \\ U^h &: d_h \times d_h \end{aligned}$$

The GRU operates using a reset gate and an update gate whereas the LSTM has an input, forget, update and output gate. For the GRU, the reset gate sits between the previous activation and the next candidate activation to forget previous state, and the update gate decides how much of the candidate activation to use in updating the cell state. GRUs have fewer parameters and thus train a bit faster or need less data to generalize hence it is computationally more efficient due to the less complex structure. The GRU unit controls the information flow like the LSTM but without the memory unit, it just exposes the full hidden content without any control.

Defensive distillation to prevent attacks that are based on gradients. The idea is to use softmax with a temperature  $T$ . Train a teacher network at  $T \gg 1$  and collect prediction probabilities  $p(x)$  for all the training samples  $x$ . Train a second network at the same large  $T \gg 1$  using  $p(x)$  as soft labels instead of 0/1 hard labels from the original dataset. At test time, use a new second network with  $T = 1$ . This would result in most predictions to have probabilities near 0/1, thus gradients are almost zero and numerically instable. This would cause attacks that are dependent on gradients to fail. Also, with a soft label provided to the second model. This uncertainty acts as an additional filter with an element of randomness to gaining a perfect match. This results in a far more robust algorithm that can spot spoofing attempts easier. Hence it adds flexibility to the algorithm's classification process so the model is less susceptible to exploitation. It creates a model whose surface is smoothed in the directions an adversary will typically try to exploit, making it difficult for them to discover adversarial input tweaks that lead to misclassification.

train the neural network with adversarial examples. Adversarial training: it is a brute force solution where we generate a lot of adversarial examples and explicitly train the model to not be fooled by each of them.

boundary attacks against black boxes can make defensive distillation fail. The idea is to move along decision boundary to the predicted class of the target image but always in the wrong way, use sample perturbations from a distribution to get a synth image that looks close to target image but has wrong prediction, only accept if the perturbed image is adversarial. This approach has iteration of two steps to get close to the target image, This can prevent defensive distillation because methods that defend by hide/shatter gradients are ineffective against the boundary attack.

if the  $g_w$  is not capped. when we try to optimize the objective function, one can minimize the object by making  $g_w$  very negative and not bothering about the first term. Hence the goals encoded by the losses  $f_{\pi}$  would not be achieved.

if  $g_w$  is capped, it yields a better result as this means that one can not minimize the objective by making  $g_w$  very small. This will hence result in the goals encoded by the losses  $f_w$  and  $q_w$  to be achieved when we optimize this.

$$\begin{aligned} P(00) &= 0.8 \times 0.8 = 0.64 \\ P(01) &= 0.8 \times 0.8 = 0.64 \\ P(10) &= 0.8 \times 0.8 = 0.64 \\ P(11) &= 0.8 \times 0.8 = 0.64 \end{aligned}$$